

```
1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 from sklearn.model_selection import train_test_split
6 from sklearn.svm import SVC
7 from sklearn.metrics import accuracy_score
8

1 data = pd.read_csv("machine_failure.csv")
2 data
```

|      | UDI   | Product ID | Type | Air temperature [K] | Process temperature [K] | Rotational speed [rpm] | Torque [Nm] | Tool wear [min] | Machine failure |
|------|-------|------------|------|---------------------|-------------------------|------------------------|-------------|-----------------|-----------------|
| 0    | 1     | M14860     | M    | 298.1               | 308.6                   | 1551                   | 42.8        | 0               |                 |
| 1    | 2     | L47181     | L    | 298.2               | 308.7                   | 1408                   | 46.3        | 3               |                 |
| 2    | 3     | L47182     | L    | 298.1               | 308.5                   | 1498                   | 49.4        | 5               |                 |
| 3    | 4     | L47183     | L    | 298.2               | 308.6                   | 1433                   | 39.5        | 7               |                 |
| 4    | 5     | L47184     | L    | 298.2               | 308.7                   | 1408                   | 40.0        | 9               |                 |
| ...  | ...   | ...        | ...  | ...                 | ...                     | ...                    | ...         | ...             |                 |
| 9995 | 9996  | M24855     | M    | 298.8               | 308.4                   | 1604                   | 29.5        | 14              |                 |
| 9996 | 9997  | H39410     | H    | 298.9               | 308.4                   | 1632                   | 31.8        | 17              |                 |
| 9997 | 9998  | M24857     | M    | 299.0               | 308.6                   | 1645                   | 33.4        | 22              |                 |
| 9998 | 9999  | H39412     | H    | 299.0               | 308.7                   | 1408                   | 48.5        | 25              |                 |
| 9999 | 10000 | M24859     | M    | 299.0               | 308.7                   | 1500                   | 40.2        | 30              |                 |

Data Preprocessing

```
1 data.columns

Index(['UDI', 'Product ID', 'Type', 'Air temperature [K]',
      'Process temperature [K]', 'Rotational speed [rpm]', 'Torque [Nm]',
      'Tool wear [min]', 'Machine failure', 'TWF', 'HDF', 'PWF', 'OSF',
      'RNF'],
      dtype='object')
```

```
1 data.nunique()

UDI                10000
Product ID         10000
Type                 3
Air temperature [K]    93
Process temperature [K]  82
Rotational speed [rpm] 941
Torque [Nm]          577
Tool wear [min]       246
Machine failure        2
TWF                   2
HDF                   2
PWF                   2
OSF                   2
RNF                   2
dtype: int64
```

```
1 data.isnull().sum()

UDI                0
Product ID         0
Type               0
Air temperature [K] 0
Process temperature [K] 0
Rotational speed [rpm] 0
Torque [Nm]        0
Tool wear [min]    0
Machine failure    0
TWF                0
HDF                0
PWF                0
OSF                0
RNF                0
dtype: int64
```

```
1 data.rename(columns= {'TWF': 'Tool Wear Failure','HDF' : 'Head Dissipation Failure', 'PWF' : 'Power Failure', 'OSF' : 'Overstrain Fail
2 data
```

|      | UDI   | Product ID | Type | Air temperature [K] | Process temperature [K] | Rotational speed [rpm] | Torque [Nm] | Tool wear [min] | Machine failure |
|------|-------|------------|------|---------------------|-------------------------|------------------------|-------------|-----------------|-----------------|
| 0    | 1     | M14860     | M    | 298.1               | 308.6                   | 1551                   | 42.8        | 0               |                 |
| 1    | 2     | L47181     | L    | 298.2               | 308.7                   | 1408                   | 46.3        | 3               |                 |
| 2    | 3     | L47182     | L    | 298.1               | 308.5                   | 1498                   | 49.4        | 5               |                 |
| 3    | 4     | L47183     | L    | 298.2               | 308.6                   | 1433                   | 39.5        | 7               |                 |
| 4    | 5     | L47184     | L    | 298.2               | 308.7                   | 1408                   | 40.0        | 9               |                 |
| ...  | ...   | ...        | ...  | ...                 | ...                     | ...                    | ...         | ...             | ...             |
| 9995 | 9996  | M24855     | M    | 298.8               | 308.4                   | 1604                   | 29.5        | 14              |                 |
| 9996 | 9997  | H39410     | H    | 298.9               | 308.4                   | 1632                   | 31.8        | 17              |                 |
| 9997 | 9998  | M24857     | M    | 299.0               | 308.6                   | 1645                   | 33.4        | 22              |                 |
| 9998 | 9999  | H39412     | H    | 299.0               | 308.7                   | 1408                   | 48.5        | 25              |                 |
| 9999 | 10000 | M24859     | M    | 299.0               | 308.7                   | 1500                   | 40.2        | 30              |                 |

```
1 data = data.drop(['UDI', 'Product ID'], axis=1)
2 data.head()
```

|   | Type | Air temperature [K] | Process temperature [K] | Rotational speed [rpm] | Torque [Nm] | Tool wear [min] | Machine failure | Tool Wear Failure | Dis |
|---|------|---------------------|-------------------------|------------------------|-------------|-----------------|-----------------|-------------------|-----|
| 0 | M    | 298.1               | 308.6                   | 1551                   | 42.8        | 0               | 0               | 0                 |     |
| 1 | L    | 298.2               | 308.7                   | 1408                   | 46.3        | 3               | 0               | 0                 |     |
| 2 | L    | 298.1               | 308.5                   | 1498                   | 49.4        | 5               | 0               | 0                 |     |
| 3 | L    | 298.2               | 308.6                   | 1433                   | 39.5        | 7               | 0               | 0                 |     |

```
1 data.rename(columns={'Type' : 'Quality Type'}, inplace = True)
2 data
```

|      | Quality Type | Air temperature [K] | Process temperature [K] | Rotational speed [rpm] | Torque [Nm] | Tool wear [min] | Machine failure | Tool Wear Failure | Dis |
|------|--------------|---------------------|-------------------------|------------------------|-------------|-----------------|-----------------|-------------------|-----|
| 0    | M            | 298.1               | 308.6                   | 1551                   | 42.8        | 0               | 0               | 0                 |     |
| 1    | L            | 298.2               | 308.7                   | 1408                   | 46.3        | 3               | 0               | 0                 |     |
| 2    | L            | 298.1               | 308.5                   | 1498                   | 49.4        | 5               | 0               | 0                 |     |
| 3    | L            | 298.2               | 308.6                   | 1433                   | 39.5        | 7               | 0               | 0                 |     |
| 4    | L            | 298.2               | 308.7                   | 1408                   | 40.0        | 9               | 0               | 0                 |     |
| ...  | ...          | ...                 | ...                     | ...                    | ...         | ...             | ...             | ...               | ... |
| 9995 | M            | 298.8               | 308.4                   | 1604                   | 29.5        | 14              | 0               | 0                 |     |
| 9996 | H            | 298.9               | 308.4                   | 1632                   | 31.8        | 17              | 0               | 0                 |     |
| 9997 | M            | 299.0               | 308.6                   | 1645                   | 33.4        | 22              | 0               | 0                 |     |
| 9998 | H            | 299.0               | 308.7                   | 1408                   | 48.5        | 25              | 0               | 0                 |     |
| 9999 | M            | 299.0               | 308.7                   | 1500                   | 40.2        | 30              | 0               | 0                 |     |

```
1 qual_map = {'L' : 'Low', 'M' : 'Medium', 'H' : 'High'}
2 data['Quality Type'] = data['Quality Type'].map(qual_map)
3 data
```

|   | Quality Type | Air temperature [K] | Process temperature [K] | Rotational speed [rpm] | Torque [Nm] | Tool wear [min] | Machine failure | Tool wear failure |
|---|--------------|---------------------|-------------------------|------------------------|-------------|-----------------|-----------------|-------------------|
| 0 | Medium       | 298.1               | 308.6                   | 1551                   | 42.8        | 0               | 0               |                   |
| 1 | Low          | 298.2               | 308.7                   | 1408                   | 46.3        | 3               | 0               |                   |
| 2 | Low          | 298.1               | 308.5                   | 1498                   | 49.4        | 5               | 0               |                   |
| 3 | Low          | 298.2               | 308.6                   | 1433                   | 39.5        | 7               | 0               |                   |
| 4 | Low          | 298.2               | 308.7                   | 1408                   | 40.0        | 9               | 0               |                   |

```
1 qual_map = {'Low' : -1, 'Medium' : 0, 'High' : 1}
2 data['Quality_Binary'] = data['Quality Type'].map(qual_map)
3 data
```

|      | Quality Type | Air temperature [K] | Process temperature [K] | Rotational speed [rpm] | Torque [Nm] | Tool wear [min] | Machine failure | Tool wear failure |
|------|--------------|---------------------|-------------------------|------------------------|-------------|-----------------|-----------------|-------------------|
| 0    | Medium       | 298.1               | 308.6                   | 1551                   | 42.8        | 0               | 0               |                   |
| 1    | Low          | 298.2               | 308.7                   | 1408                   | 46.3        | 3               | 0               |                   |
| 2    | Low          | 298.1               | 308.5                   | 1498                   | 49.4        | 5               | 0               |                   |
| 3    | Low          | 298.2               | 308.6                   | 1433                   | 39.5        | 7               | 0               |                   |
| 4    | Low          | 298.2               | 308.7                   | 1408                   | 40.0        | 9               | 0               |                   |
| ...  | ...          | ...                 | ...                     | ...                    | ...         | ...             | ...             |                   |
| 9995 | Medium       | 298.8               | 308.4                   | 1604                   | 29.5        | 14              | 0               |                   |
| 9996 | High         | 298.9               | 308.4                   | 1632                   | 31.8        | 17              | 0               |                   |
| 9997 | Medium       | 299.0               | 308.6                   | 1645                   | 33.4        | 22              | 0               |                   |
| 9998 | High         | 299.0               | 308.7                   | 1408                   | 48.5        | 25              | 0               |                   |
| 9999 | Medium       | 299.0               | 308.7                   | 1500                   | 40.2        | 30              | 0               |                   |

```
1 data.tail()
```

|      | Quality Type | Air temperature [K] | Process temperature [K] | Rotational speed [rpm] | Torque [Nm] | Tool wear [min] | Machine failure | Tool wear failure |
|------|--------------|---------------------|-------------------------|------------------------|-------------|-----------------|-----------------|-------------------|
| 9995 | Medium       | 298.8               | 308.4                   | 1604                   | 29.5        | 14              | 0               |                   |
| 9996 | High         | 298.9               | 308.4                   | 1632                   | 31.8        | 17              | 0               |                   |
| 9997 | Medium       | 299.0               | 308.6                   | 1645                   | 33.4        | 22              | 0               |                   |
| 9998 | High         | 299.0               | 308.7                   | 1408                   | 48.5        | 25              | 0               |                   |

```
1 data.describe()
```

|       | Air temperature [K] | Process temperature [K] | Rotational speed [rpm] | Torque [Nm]  | Tool wear [min] | Machine failure |
|-------|---------------------|-------------------------|------------------------|--------------|-----------------|-----------------|
| count | 10000.000000        | 10000.000000            | 10000.000000           | 10000.000000 | 10000.000000    | 10000.000000    |
| mean  | 300.004930          | 310.005560              | 1538.776100            | 39.986910    | 107.951000      | 0.000000        |
| std   | 2.000259            | 1.483734                | 179.284096             | 9.968934     | 63.654147       | 0.111111        |
| min   | 295.300000          | 305.700000              | 1168.000000            | 3.800000     | 0.000000        | 0.000000        |
| 25%   | 298.300000          | 308.800000              | 1423.000000            | 33.200000    | 53.000000       | 0.000000        |
| 50%   | 300.100000          | 310.100000              | 1503.000000            | 40.100000    | 108.000000      | 0.000000        |
| 75%   | 301.500000          | 311.100000              | 1612.000000            | 46.800000    | 162.000000      | 0.000000        |

```
1 data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 13 columns):
#   Column              Non-Null Count  Dtype
---  -
0   Quality Type        10000 non-null  object
```

```

1  Air temperature [K]      10000 non-null float64
2  Process temperature [K]  10000 non-null float64
3  Rotational speed [rpm]   10000 non-null int64
4  Torque [Nm]              10000 non-null float64
5  Tool wear [min]          10000 non-null int64
6  Machine failure          10000 non-null int64
7  Tool Wear Failure        10000 non-null int64
8  Head Dissipation Failure 10000 non-null int64
9  Power Failure            10000 non-null int64
10 Overstrain Failure        10000 non-null int64
11 Random Failures          10000 non-null int64
12 Quality_Binary           10000 non-null int64
dtypes: float64(3), int64(9), object(1)
memory usage: 1015.8+ KB

```

```
1 data.shape
```


```
(10000, 13)
```

## Data Analysis

```

1 #Outlier Detection and removal
2 data = pd.DataFrame(data)
3 numerical_cols = data.select_dtypes(include=['float64']).columns.tolist() + ['Rotational speed [rpm]']
4 # numerical_cols.append('Rotational speed [rpm]')
5 print(numerical_cols)
6
7 # Create a custom color palette for boxplots
8 colors = ['#1f77b4', '#ff7f0e', '#2ca02c', '#d62728', '#9467bd', '#8c564b', '#e377c2', '#7f7f7f', '#bcbd22', '#17becf']
9
10
11
12 fig, axes = plt.subplots(nrows = 1, ncols = len(numerical_cols), figsize = (15,5))
13
14 #Loop through nnumerical_cols and create boxplots
15 for i, column in enumerate(numerical_cols):
16     ax = axes[i]
17
18     # Customize boxplot appearance
19     boxprops = dict(linewidth=2, color=colors[i])
20     flierprops = dict(marker='o', markersize=6, markerfacecolor=colors[i], markeredgecolor='none')
21
22     ax.boxplot(data[column], boxprops=boxprops, flierprops=flierprops)
23
24
25     ax.boxplot(data[column])
26     ax.set_title(f'Boxplot for {column}')
27     ax.set_xticks([])
28
29 plt.tight_layout()
30 plt.show()

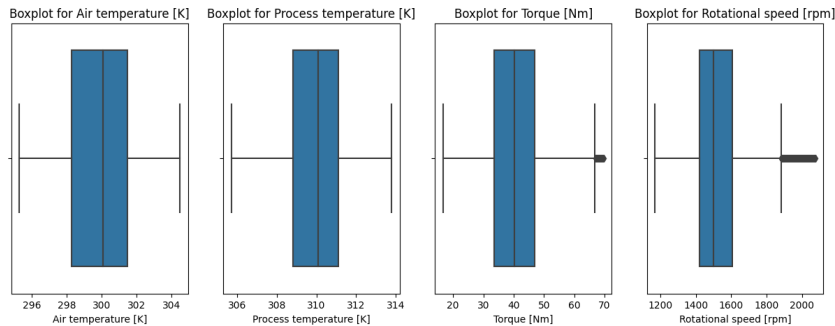
```



```
1 #Detecting outliers
2 from scipy import stats
3 # Define the z-score threshold (e.g., 3)
4 z_score_threshold = 3
5
6 # Create a mask to identify outliers for each numerical column
7 outlier_mask = np.abs(stats.zscore(data[numerical_cols])) > z_score_threshold
8
9 # Apply the mask to remove outliers from the DataFrame
10 data_no_outliers = data[~outlier_mask.any(axis=1)]
11
12 # Print the shape of the DataFrame before and after removing outliers
13 print("Original Data Shape:", data.shape)
14 print("Data Shape After Removing Outliers:", data_no_outliers.shape)
15
16 # Optionally, you can reset the index if needed
17 data_no_outliers.reset_index(drop=True, inplace=True)
```

Original Data Shape: (10000, 13)  
Data Shape After Removing Outliers: (9822, 13)

```
1 # Optionally, you can reset the index if needed
2 # data_no_outliers.reset_index(drop=True, inplace=True)
3
4 # Plot boxplots for numerical columns in data_no_outliers
5 fig, axes = plt.subplots(nrows=1, ncols=len(numerical_cols), figsize=(15, 5))
6
7 for i, column in enumerate(numerical_cols):
8     sns.boxplot(x=data_no_outliers[column], ax=axes[i])
9     axes[i].set_title(f'Boxplot for {column}')
10
```



```
1
1 data_no_outliers.describe()
```

|       | Air temperature [K] | Process temperature [K] | Rotational speed [rpm] | Torque [Nm] | Tool wear [min] | Machine failure |
|-------|---------------------|-------------------------|------------------------|-------------|-----------------|-----------------|
| count | 9822.000000         | 9822.000000             | 9822.000000            | 9822.000000 | 9822.000000     | 9822.000000     |
| mean  | 300.001161          | 310.002861              | 1526.222765            | 40.350713   | 107.912441      | 0.029831        |
| std   | 1.998035            | 1.482233                | 147.382767             | 9.448740    | 63.616929       | 0.170130        |
| min   | 295.300000          | 305.700000              | 1168.000000            | 16.700000   | 0.000000        | 0.000000        |
| 25%   | 298.300000          | 308.800000              | 1422.000000            | 33.600000   | 53.000000       | 0.000000        |
| 50%   | 300.100000          | 310.100000              | 1500.000000            | 40.200000   | 108.000000      | 0.000000        |
| 75%   | 301.500000          | 311.100000              | 1606.000000            | 46.900000   | 162.000000      | 0.000000        |

```
1 data = data_no_outliers
```

1 data

|      | Quality Type | Air temperature [K] | Process temperature [K] | Rotational speed [rpm] | Torque [Nm] | Tool wear [min] | Machine failure | Tool Wear Failure |
|------|--------------|---------------------|-------------------------|------------------------|-------------|-----------------|-----------------|-------------------|
| 0    | Medium       | 298.1               | 308.6                   | 1551                   | 42.8        | 0               | 0               |                   |
| 1    | Low          | 298.2               | 308.7                   | 1408                   | 46.3        | 3               | 0               |                   |
| 2    | Low          | 298.1               | 308.5                   | 1498                   | 49.4        | 5               | 0               |                   |
| 3    | Low          | 298.2               | 308.6                   | 1433                   | 39.5        | 7               | 0               |                   |
| 4    | Low          | 298.2               | 308.7                   | 1408                   | 40.0        | 9               | 0               |                   |
| ...  | ...          | ...                 | ...                     | ...                    | ...         | ...             | ...             |                   |
| 9817 | Medium       | 298.8               | 308.4                   | 1604                   | 29.5        | 14              | 0               |                   |
| 9818 | High         | 298.9               | 308.4                   | 1632                   | 31.8        | 17              | 0               |                   |
| 9819 | Medium       | 299.0               | 308.6                   | 1645                   | 33.4        | 22              | 0               |                   |
| 9820 | High         | 299.0               | 308.7                   | 1408                   | 48.5        | 25              | 0               |                   |
| 9821 | Medium       | 299.0               | 308.7                   | 1500                   | 40.2        | 30              | 0               |                   |

Feature Importance in the Dataset

1 data.columns

```
Index(['Quality Type', 'Air temperature [K]', 'Process temperature [K]',
      'Rotational speed [rpm]', 'Torque [Nm]', 'Tool wear [min]',
      'Machine failure', 'Tool Wear Failure', 'Head Dissipation Failure',
      'Power Failure', 'Overstrain Failure', 'Random Failures',
      'Quality_Binary'],
      dtype='object')
```

```
1 from sklearn.ensemble import RandomForestClassifier
2 from sklearn.preprocessing import LabelEncoder
3 X = data.drop(['Machine failure', 'Quality Type'], axis=1)
4 y = data['Machine failure']
```

1 X.columns

```
Index(['Air temperature [K]', 'Process temperature [K]',
      'Rotational speed [rpm]', 'Torque [Nm]', 'Tool wear [min]',
      'Tool Wear Failure', 'Head Dissipation Failure', 'Power Failure',
      'Overstrain Failure', 'Random Failures', 'Quality_Binary'],
      dtype='object')
```

1 X.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9822 entries, 0 to 9821
Data columns (total 11 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Air temperature [K]                   9822 non-null   float64
1   Process temperature [K]               9822 non-null   float64
2   Rotational speed [rpm]                9822 non-null   int64
3   Torque [Nm]                           9822 non-null   float64
4   Tool wear [min]                       9822 non-null   int64
5   Tool Wear Failure                     9822 non-null   int64
6   Head Dissipation Failure              9822 non-null   int64
7   Power Failure                         9822 non-null   int64
8   Overstrain Failure                    9822 non-null   int64
9   Random Failures                       9822 non-null   int64
10  Quality_Binary                        9822 non-null   int64
dtypes: float64(3), int64(8)
memory usage: 844.2 KB
```

```
1 label_encoder = LabelEncoder()
2 y = label_encoder.fit_transform(y)
```

1 y

```
array([0, 0, 0, ..., 0, 0, 0], dtype=int64)
```

```
1 model = RandomForestClassifier()  
2 model
```

▼ RandomForestClassifier

RandomForestClassifier()

```
1 model.fit(X, y)
```

▼ RandomForestClassifier

RandomForestClassifier()

```
1 #Get feature importance from the trained mmodel  
2 feature_importance = model.feature_importances_  
3 feature_importance  
  
array([2.07506225e-02, 1.43808631e-02, 3.75407208e-02, 7.51327823e-02,  
       4.66449261e-02, 1.29783416e-01, 3.42184822e-01, 1.01555714e-01,  
       2.27190008e-01, 3.31725977e-04, 4.50439814e-03])
```


```
1 # Create a DataFrame to display feature importances  
2 importance_df = pd.DataFrame({'Feature': X.columns, 'Importance': feature_importance})  
3 importance_df
```

|    | Feature                  | Importance |
|----|--------------------------|------------|
| 0  | Air temperature [K]      | 0.020751   |
| 1  | Process temperature [K]  | 0.014381   |
| 2  | Rotational speed [rpm]   | 0.037541   |
| 3  | Torque [Nm]              | 0.075133   |
| 4  | Tool wear [min]          | 0.046645   |
| 5  | Tool Wear Failure        | 0.129783   |
| 6  | Head Dissipation Failure | 0.342185   |
| 7  | Power Failure            | 0.101556   |
| 8  | Overstrain Failure       | 0.227190   |
| 9  | Random Failures          | 0.000332   |
| 10 | Quality_Binary           | 0.004504   |

```
1 # Sort the DataFrame by importance values in descending order  
2 importance_df = importance_df.sort_values(by='Importance', ascending=False)  
3 importance_df
```

|    | Feature                  | Importance |
|----|--------------------------|------------|
| 6  | Head Dissipation Failure | 0.342185   |
| 8  | Overstrain Failure       | 0.227190   |
| 5  | Tool Wear Failure        | 0.129783   |
| 7  | Power Failure            | 0.101556   |
| 3  | Torque [Nm]              | 0.075133   |
| 4  | Tool wear [min]          | 0.046645   |
| 2  | Rotational speed [rpm]   | 0.037541   |
| 0  | Air temperature [K]      | 0.020751   |
| 1  | Process temperature [K]  | 0.014381   |
| 10 | Quality_Binary           | 0.004504   |
| 9  | Random Failures          | 0.000332   |

```
1 # Create a bar plot to visualize feature importances  
2 plt.figure(figsize=(10, 6))  
3 plt.barh(importance_df['Feature'], importance_df['Importance'])  
4 plt.xlabel('Feature Importance')  
5 plt.ylabel('Feature')  
6 plt.title('Feature Importance Scores')  
7 plt.gca().invert_yaxis() # Invert the y-axis to display the most important features at the top  
8 plt.show()
```



```
1 data.columns

Index(['Quality Type', 'Air temperature [K]', 'Process temperature [K]',
      'Rotational speed [rpm]', 'Torque [Nm]', 'Tool wear [min]',
      'Machine failure', 'Tool Wear Failure', 'Head Dissipation Failure',
      'Power Failure', 'Overstrain Failure', 'Random Failures',
      'Quality_Binary'],
      dtype='object')

1 # Sort importance_df to get the top 8 important features
2 top_features = importance_df.nlargest(8, 'Importance')['Feature']
3
4 # Create a new DataFrame with the top 8 features and the target variable 'Machine Failures'
5 new_data = data[list(top_features) + ['Machine failure']] + ['Quality_Binary']]
6 new_data
```

|      | Head<br>Dissipation<br>Failure | Overstrain<br>Failure | Tool<br>Wear<br>Failure | Power<br>Failure | Torque<br>[Nm] | Tool<br>wear<br>[min] | Rotational<br>speed<br>[rpm] | temperi |
|------|--------------------------------|-----------------------|-------------------------|------------------|----------------|-----------------------|------------------------------|---------|
| 0    | 0                              | 0                     | 0                       | 0                | 42.8           | 0                     | 1551                         |         |
| 1    | 0                              | 0                     | 0                       | 0                | 46.3           | 3                     | 1408                         |         |
| 2    | 0                              | 0                     | 0                       | 0                | 49.4           | 5                     | 1498                         |         |
| 3    | 0                              | 0                     | 0                       | 0                | 39.5           | 7                     | 1433                         |         |
| 4    | 0                              | 0                     | 0                       | 0                | 40.0           | 9                     | 1408                         |         |
| ...  | ...                            | ...                   | ...                     | ...              | ...            | ...                   | ...                          |         |
| 9817 | 0                              | 0                     | 0                       | 0                | 29.5           | 14                    | 1604                         |         |
| 9818 | 0                              | 0                     | 0                       | 0                | 31.8           | 17                    | 1632                         |         |
| 9819 | 0                              | 0                     | 0                       | 0                | 33.4           | 22                    | 1645                         |         |
| 9820 | 0                              | 0                     | 0                       | 0                | 48.5           | 25                    | 1408                         |         |
| 9821 | 0                              | 0                     | 0                       | 0                | 40.2           | 30                    | 1500                         |         |

```
1 len(data.columns)

13

1 len(new_data.columns)

10

1 # Get the set of columns in the original data DataFrame
2 data_columns = set(data.columns)
3
```



```

4 # Get the set of columns in the new_data DataFrame
5 new_data_columns = set(new_data.columns)

1 # Find the columns that are in 'data' but not in 'new_data'
2 columns_not_in_new_data = data_columns - new_data_columns
3
4 # Now, 'columns_not_in_new_data' contains the columns that are in 'data' but not in 'new_data'
5 print("Columns not in new_data:", columns_not_in_new_data)

Columns not in new_data: {'Random Failures', 'Quality Type', 'Process temperature [K]'}

```

### Multivariate Analysis

- Multivariate Analysis involves exploring relationships between multiple variables simultaneously. Steps -
- Standardize the data
- Apply PCA
- Select number of Components
- Transform Data
- Visualization


```

1 from sklearn.preprocessing import StandardScaler
2 from sklearn.decomposition import PCA

1 # Standardize the data
2 scaler = StandardScaler()
3 scaled_data = scaler.fit_transform(new_data.drop('Machine failure', axis=1))

1 pca = PCA()
2 pca.fit(scaled_data)
3 explained_variance = pca.explained_variance_ratio_
4 cumulative_variance = np.cumsum(explained_variance)
5 # Plot explained variance to decide the number of components
6 plt.plot(cumulative_variance)
7 plt.xlabel('Number of Components')
8 plt.ylabel('Cumulative Explained Variance')
9 plt.grid()
10 plt.show()

```



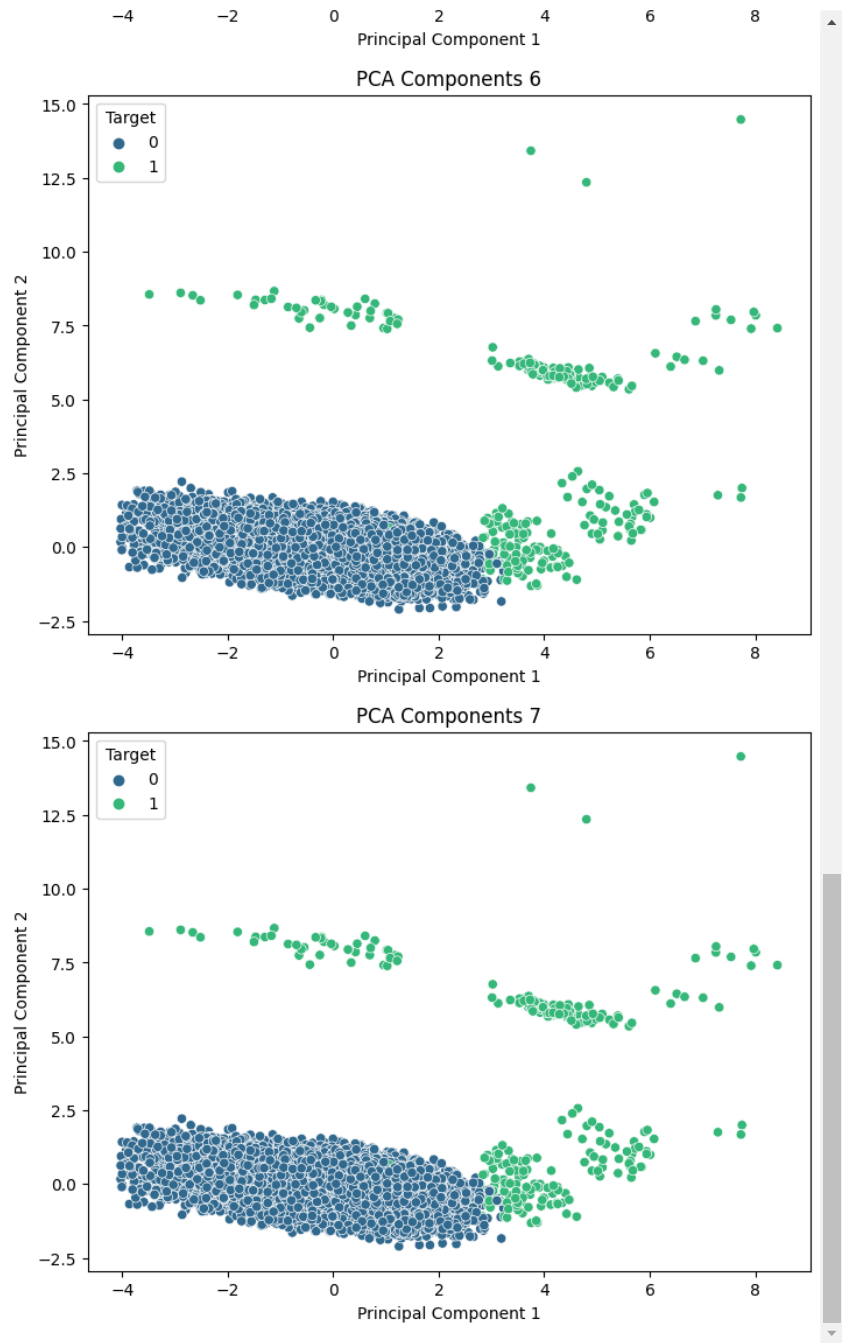
```

1 from sklearn.decomposition import PCA
2
3 # Assuming you have a dataset named 'scaled_data' and 'y' as your target variable
4
5 # Create a function to visualize reduced dimensions
6 def visualize_pca_components(components, labels):
7     pca = PCA(n_components=components)
8     pca_result = pca.fit_transform(scaled_data)
9
10    # Create a scatter plot of the first two principal components
11    plt.figure(figsize=(8, 6))
12    sns.scatterplot(x=pca_result[:, 0], y=pca_result[:, 1], hue=labels, palette="viridis")
13    plt.title(f'PCA Components {components}')
14    plt.xlabel('Principal Component 1')
15    plt.ylabel('Principal Component 2')

```

```
16 plt.legend(title='Target', loc='best')
17 plt.show()
18
19
```

```
1 # Call the function with different numbers of components
2 visualize_pca_components(2, y) # Visualize with 2 components
3 visualize_pca_components(3, y) # Visualize with 3 components
4 visualize_pca_components(4, y) # Visualize with 4 components
5 visualize_pca_components(5, y)
6 visualize_pca_components(6, y)
7 visualize_pca_components(7, y)
```



```

1 # Apply PCA
2 pca = PCA(n_components=7) # Choose the number of components you want to retain
3 pca_result = pca.fit_transform(scaled_data)
4 pca_result

array([[ -0.06960303, -1.43007683, -0.70784679, ...,  0.58585201,
         0.31328833,  0.64121409],
       [ 0.83414652, -1.36355164, -0.58063326, ..., -0.38344926,
         0.91343419,  0.42246668 ],
       [ 0.65875132, -1.2900605 , -0.60034887, ..., -0.30771422,
         0.88917705,  0.48714783],
       ...,
       [-1.13026506, -0.93171705, -0.28549213, ...,  0.76713283,
         0.13257012,  0.51598165],
       [ 0.94764774, -1.57035252, -0.71710542, ...,  1.0566524 ,
        -0.33378582,  0.16785075],
       [-0.01100348, -1.09099993, -0.42095047, ...,  0.48830018,
         0.1484816 ,  0.30929213]])

```

### Applying KNN

```

1 from sklearn.neighbors import KNeighborsClassifier
2 from sklearn.metrics import accuracy_score

1 # Split the data into training and testing sets
2 X_train, X_test, y_train, y_test = train_test_split(pca_result, y, test_size=0.2, random_state=42)

1 # Feature scaling (standardization)
2 scaler = StandardScaler()
3 X_train_scaled = scaler.fit_transform(X_train)
4 X_test_scaled = scaler.transform(X_test)

1 X_train_scaled.shape

(7857, 7)


```

### Search for optimal number of Neighbors

```

1 from sklearn.model_selection import GridSearchCV
2 # Initialize the KNN classifier
3 knn_classifier = KNeighborsClassifier()
4 # Define a range of 'k' values to search
5 param_grid = {'n_neighbors': [1, 3, 5, 7, 9]}
6 # Perform grid search with cross-validation
7 grid_search = GridSearchCV(knn_classifier, param_grid, cv=5)
8 grid_search.fit(X_train_scaled, y_train)

```



```
1 # Get the best 'k' value from the grid search
2 best_k = grid_search.best_params_['n_neighbors']
3 best_k
```

```
3
```

```
1 # Initialize and train the KNN classifier (you can choose the number of neighbors 'n_neighbors')
2 knn_classifier = KNeighborsClassifier(n_neighbors=3)
3 knn_classifier
```

```
▼ KNeighborsClassifier
KNeighborsClassifier(n_neighbors=3)
```

```
1 # Train the classifier on your data
2 knn_classifier.fit(X_train_scaled, y_train)
```

```
▼ KNeighborsClassifier
KNeighborsClassifier(n_neighbors=3)
```

```
1 # Make predictions on test data
2 y_pred = knn_classifier.predict(X_test_scaled)
```

```
1 # Evaluate the classifier's performance
2 accuracy = accuracy_score(y_test, y_pred)
3 print(f'Accuracy with k=3: {accuracy}')
```

```
Accuracy with k=3: 0.9994910941475827
```

### Checking other Accuracy measures and Scores

```
1 from sklearn.metrics import confusion_matrix, precision_score, recall_score, f1_score, roc_curve, roc_auc_score
2 from sklearn.model_selection import cross_val_score
```

```
1 # Calculate the confusion matrix
2 conf_matrix = confusion_matrix(y_test, y_pred)
3 print('Confusion Matrix:')
4 print(conf_matrix)
```

```
Confusion Matrix:
[[1905   0]
 [   1  59]]
```

```
1 # Calculate precision, recall, and F1-score
2 precision = precision_score(y_test, y_pred)
3 recall = recall_score(y_test, y_pred)
4 f1 = f1_score(y_test, y_pred)
5 print(f'Precision: {precision:.4f}')
6 print(f'Recall: {recall:.4f}')
7 print(f'F1-Score: {f1:.4f}')
```


```
Precision: 1.0000
Recall: 0.9833
F1-Score: 0.9916
```

```
1 # Calculate ROC curve and AUC
2 y_prob = knn_classifier.predict_proba(X_test_scaled)[:, 1]
3 fpr, tpr, thresholds = roc_curve(y_test, y_prob)
4 roc_auc = roc_auc_score(y_test, y_prob)
5 print(f'ROC AUC: {roc_auc:.4f}')
6
```

```
ROC AUC: 0.9916
```

```
1 # Plot ROC curve
2 import matplotlib.pyplot as plt
3 plt.figure(figsize=(8, 6))
4 plt.plot(fpr, tpr, linewidth=2, label=f'ROC Curve (AUC = {roc_auc:.2f})')
5 plt.plot([0, 1], [0, 1], 'k--', linewidth=2)
6 plt.xlim([0.0, 1.0])
7 plt.ylim([0.0, 1.05])
8 plt.xlabel('False Positive Rate')
9 plt.ylabel('True Positive Rate')
10 plt.title('Receiver Operating Characteristic (ROC)')
```

```
11 plt.legend(loc='lower right')
12 plt.show()
```



```
1 # Perform cross-validation (5-fold in this example)
2 cv_scores = cross_val_score(knn_classifier, X_train_scaled, y_train, cv=5)
3 print(f'Cross-Validation Scores: {cv_scores}')
4 print(f'Average Cross-Validation Score: {cv_scores.mean():.4f}')
```

Cross-Validation Scores: [0.99936387 0.99745547 0.99936346 0.99936346 0.99936346]  
Average Cross-Validation Score: 0.9990

#### Applying Support-Vector Machines in this Dataset

```
1 from sklearn.svm import SVC
2
3 # Initialize the SVM classifier with the RBF kernel
4 svm_classifier = SVC(kernel='rbf', random_state=42)
5 svm_classifier
```

```
SVC
SVC(random_state=42)
```

```
1 # Fit the SVM model to the training data
2 svm_classifier.fit(X_train_scaled, y_train)
```

```
SVC
SVC(random_state=42)
```

```
1 from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
2
3 # Make predictions on the testing data
4 y_pred = svm_classifier.predict(X_test_scaled)
5
6 # Evaluate the model
7 accuracy = accuracy_score(y_test, y_pred)
8 conf_matrix = confusion_matrix(y_test, y_pred)
9 classification_rep = classification_report(y_test, y_pred)
10
11 # Print the results
12 print("Accuracy:", accuracy)
13 print("Confusion Matrix:\n", conf_matrix)
14 print("Classification Report:\n", classification_rep)
```

```
Accuracy: 0.9994910941475827
Confusion Matrix:
[[1905  0]
 [  1 59]]
```